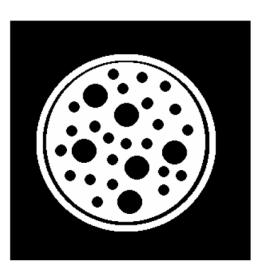
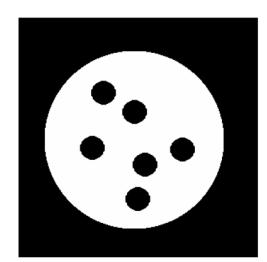


Edges and Binary Image Analysis



Notes Adapted from Prof. Kristen Grauman



Previously

- Filters allow local image neighborhood to influence our description and features
 - Smoothing to reduce noise
 - Derivatives to locate contrast, gradient

Today

- Edge detection and matching
 - process the image gradient to find curves/contours
 - comparing contours
- Binary image analysis
 - blobs and regions

Edge detection

- Goal: map image from 2d array of pixels to a set of curves or line segments or contours.
- Why?

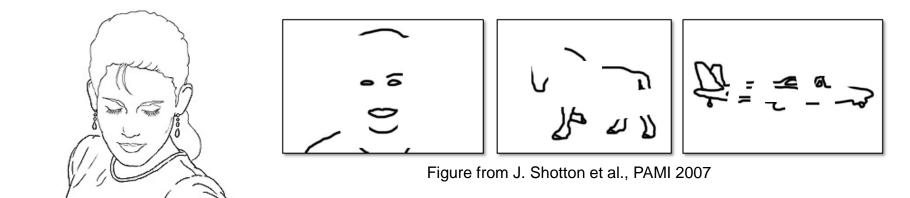


Figure from D. Lowe

Main idea: look for strong gradients, post-process



Gradients -> edges



- 1. Smoothing: suppress noise
- 2. Edge enhancement: filter for contrast
- 3. Edge localization

Determine which local maxima from filter output are actually edges vs. noise

• Threshold, Thin

Thresholding

- Choose a threshold value t
- Set any pixels less than t to zero (off)
- Set any pixels greater than or equal to t to one (on)

Original image



Gradient magnitude image



Thresholding gradient with a lower threshold



Thresholding gradient with a higher threshold



- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Non-maximum suppression:
 - Thin wide "ridges" down to single pixel width
- Linking and thresholding (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

- MATLAB: edge(image, 'canny');
- >>help edge



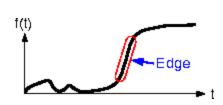
original image (Lena)

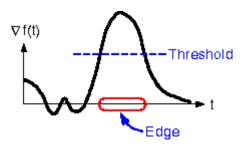


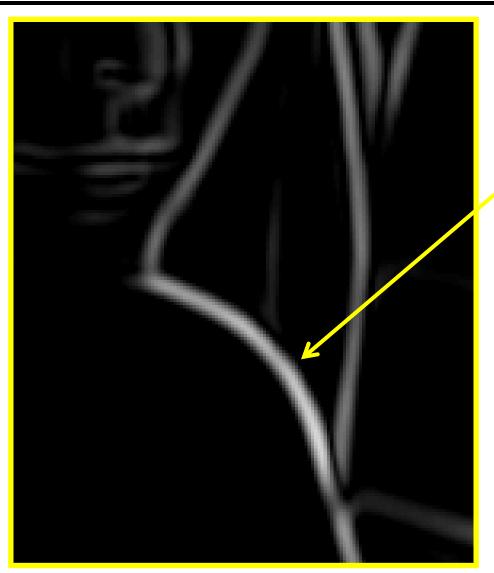
norm of the gradient



thresholding

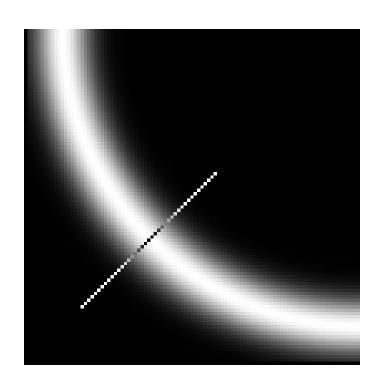


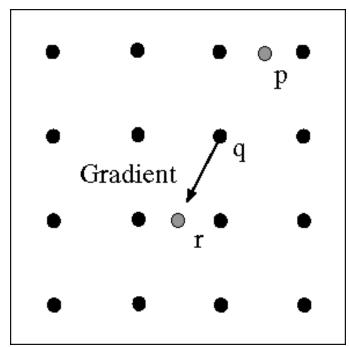




How to turn these thick regions of the gradient into curves?

Non-maximum suppression





Check if pixel is local maximum along gradient direction, select single max across width of the edge

requires checking interpolated pixels p and r

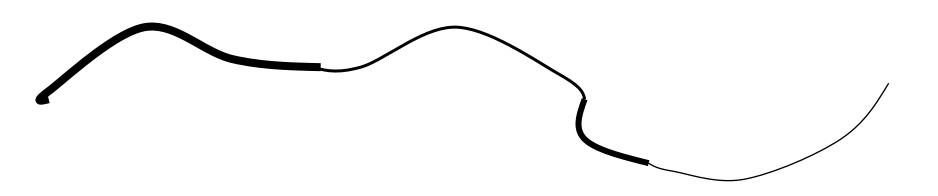


Problem:
pixels along
this edge
didn't
survive the
thresholding

thinning (non-maximum suppression)

Hysteresis thresholding

 Use a high threshold to start edge curves, and a low threshold to continue them.



Hysteresis thresholding



original image



high threshold (strong edges)



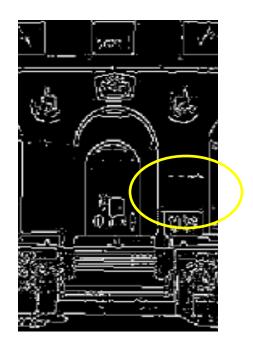
low threshold (weak edges)



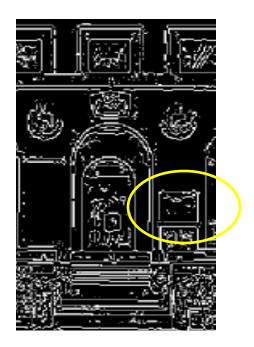
hysteresis threshold

Source: L. Fei-Fei

Hysteresis thresholding



high threshold (strong edges)



low threshold (weak edges)



hysteresis threshold

Recap: Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Non-maximum suppression:
 - Thin wide "ridges" down to single pixel width
- Linking and thresholding (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

- MATLAB: edge(image, 'canny');
- >>help edge

Low-level edges vs. perceived contours













Background

Texture

Shadows

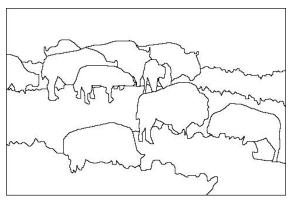
Low-level edges vs. perceived contours

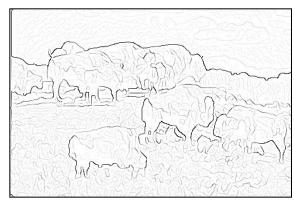
image

human segmentation

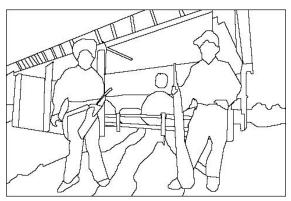
gradient magnitude











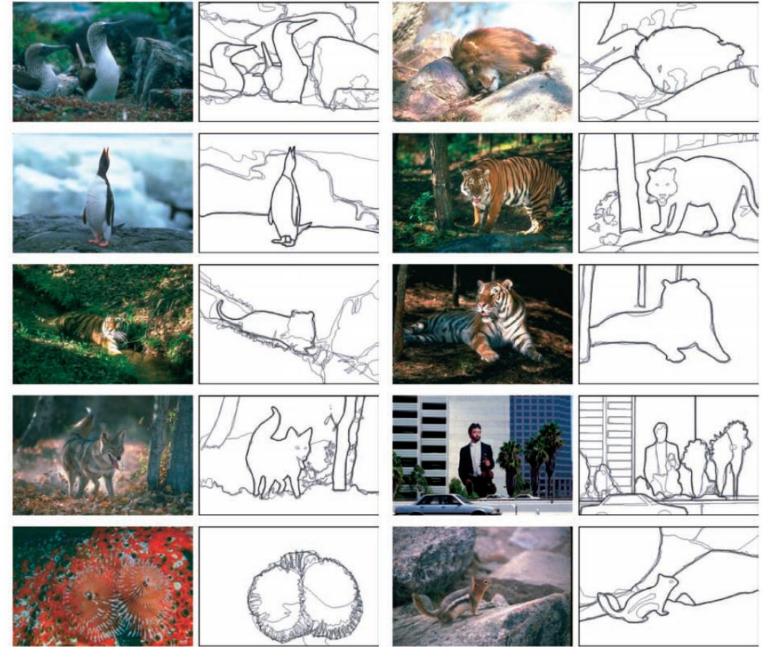


Berkeley segmentation database:

http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/

Source: L. Lazebnik

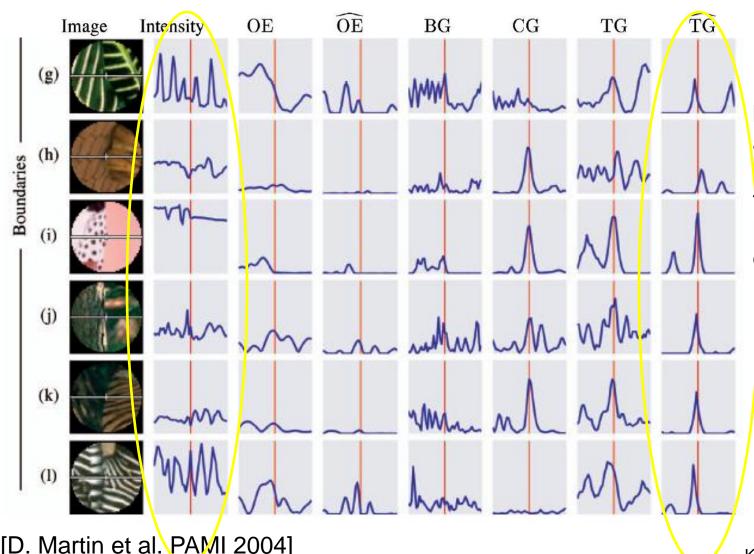
Learn from humans which combination of features is most indicative of a "good" contour?



Human-marked segment boundaries

[D. Martin et al. PAMI 2004]

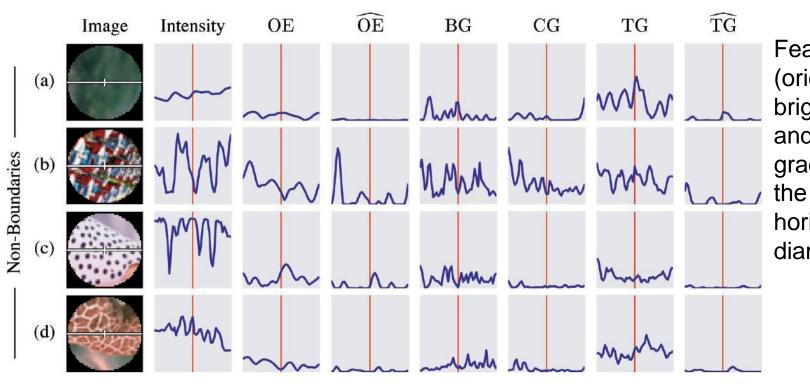
What features are responsible for perceived edges?



Feature profiles (oriented energy, brightness, color, and texture gradients) along the patch's horizontal diameter

[D. Martin et al. PAMI 2004] Kristen Grauman, UT-Austin

What features are responsible for perceived edges?



Feature profiles
(oriented energy,
brightness, color,
and texture
gradients) along
the patch's
horizontal
diameter



[D. Martin et al. PAMI 2004]

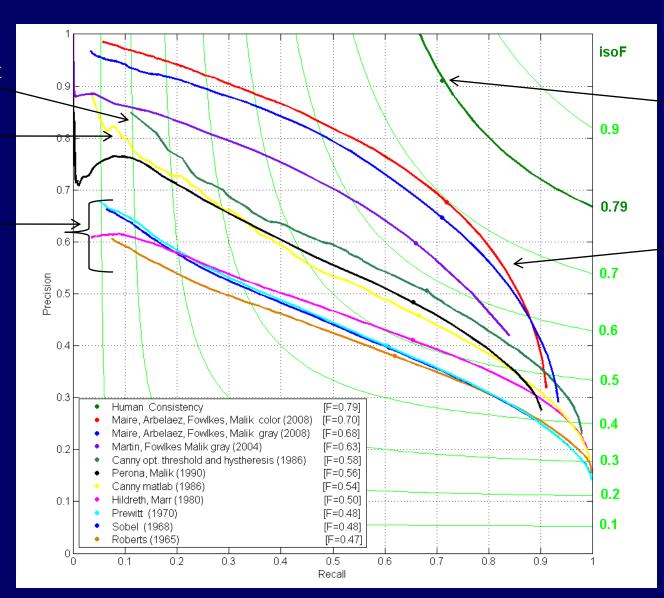
Kristen Grauman, UT-Austin

State-of-the-Art in Contour Detection

Canny+opt thresholds

Canny

Prewitt, Sobel, Roberts



Human agreement

Learned with combined features

Today

- Edge detection and matching
 - process the image gradient to find curves/contours
 - comparing contours
- Binary image analysis
 - blobs and regions

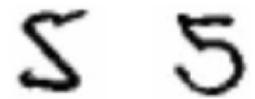


Fig. 1. Examples of two handwritten digits. In terms of pixel-to-pixel comparisons, these two images are quite different, but to the human observer, the shapes appear to be similar.

Chamfer distance

Average distance to nearest feature

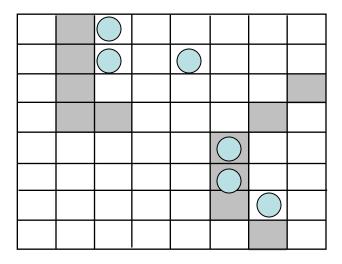
$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$

 $I={\sf Set}$ of points in image

T= Set of points on (shifted) template

 $d_I(t) =$ Minimum distance between point t and some point in I

Chamfer distance



$$D_{chamfer}(T,I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$

Chamfer distance

Average distance to nearest feature

$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$





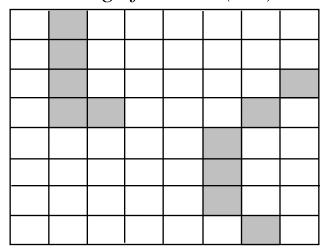
Edge image

How is the measure different than just filtering with a mask having the shape points?

How expensive is a naïve implementation?

Distance transform

Image features (2D)



Distance Transform

| 1 | 0 | 1 | 2 | 3 | 4 | 3 | 2 |
|---|---|---|---|---|---|---|---|
| 1 | 0 | 1 | 2 | 3 | 3 | 2 | 1 |
| 1 | 0 | 1 | 2 | 3 | 2 | 1 | 0 |
| 1 | 0 | 0 | 1 | 2 | 1 | 0 | 1 |
| 2 | 1 | 1 | 2 | 1 | 0 | 1 | 2 |
| 3 | 2 | 2 | 2 | 1 | 0 | 1 | 2 |
| 4 | 3 | 3 | 2 | 1 | 0 | 1 | 2 |
| 5 | 4 | 4 | 3 | 2 | 1 | 0 | 1 |

Distance Transform is a function $D(\cdot)$ that for each image pixel p assigns a non-negative number D(p) corresponding to distance from p to the nearest feature in the image I

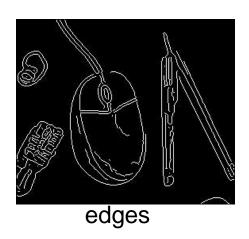
Features could be edge points, foreground points,...

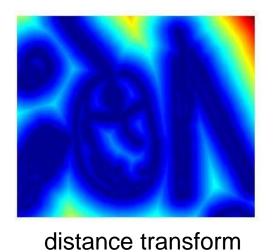
Source: Yuri Boykov

Distance transform



original





Value at (x,y) tells how far that position is from the nearest edge point (or other binary mage structure)

>> help bwdist

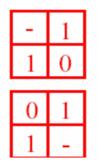
Distance transform (1D)

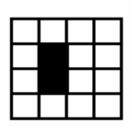
Two pass O(n) algorithm for 1D L₁ norm

```
1. Initialize: For all j
D[j] \leftarrow 1_{\mathbf{P}}[j] \qquad \text{# 0 if j is in } \mathbf{P}, \text{ infinity otherwise}
```

Distance Transform (2D)

- 2D case analogous to 1D
 - Initialization
 - Forward and backward pass
 - Fwd pass finds closest above and to left
 - Bwd pass finds closest below and to right





| 8 | 8 | 8 | 8 |
|---|---|---|---|
| 8 | 0 | 8 | ∞ |
| 8 | 0 | 8 | 8 |
| 8 | 8 | 8 | 8 |

| ∞ | 8 | 8 | 8 |
|---|---|---|---|
| ∞ | 0 | 1 | 8 |
| 8 | 0 | 8 | 8 |
| 8 | 8 | 8 | 8 |

| _∞ | 8 | 8 | 8 |
|--------------|---|---|---|
| 8 | 0 | 1 | 2 |
| 8 | 0 | 1 | 2 |
| 8 | 1 | 2 | 3 |


| 2 | 1 | 2 | 3 |
|---|---|---|---|
| 1 | 0 | 1 | 2 |
| 1 | 0 | 1 | 2 |
| 2 | 1 | 2 | 3 |

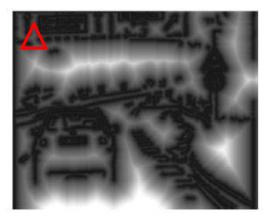
Chamfer distance

Average distance to nearest feature

$$D_{chamfer}(T,I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$







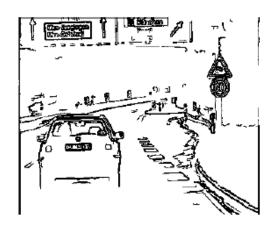
Edge image

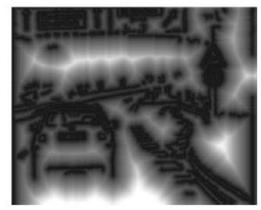
Distance transform image

Chamfer distance









Edge image

Distance transform image

Chamfer distance: properties

- Sensitive to scale and rotation
- Tolerant of small shape changes, clutter
- Need large number of template shapes
- Inexpensive way to match shapes

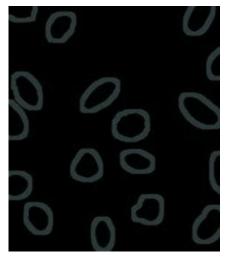
Today


- Edge detection and matching
 - process the image gradient to find curves/contours
 - comparing contours
- Binary image analysis
 - blobs and regions

Binary images

















A mall-conflict switch for the Legion estihukhape memory offers

The Decision Moral for the execution is more beef and stage of the in-contributed in give in restriction. It can provide a large of the Decision and give in Distriction and provide of the filter best and in a stage of the Distriction and provide of the configuration of the contribution of the contributio

The photo concepts is a dependent operator of these payoness, which we feel because a most dependent entering an entering an expectation of the dependent of the control of the dependent of the


Binary image analysis: basic steps

- Convert the image into binary form
 - Thresholding
- Clean up the thresholded image
 - Morphological operators
- Extract separate blobs
 - Connected components
- Describe the blobs with region properties

Binary images

- Two pixel values
 - Foreground and background
 - Mark region(s) of interest

| 1 | 1 | 0 | 1 | 1 | 1 | 0 | 1 |
|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 1 | 1 | 0 | 0 | 0 | 1 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| 1 | 1 | 1 | 1 | 0 | 1 | 0 | 1 |
| 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 0 | 1 | 0 | 0 | 0 | 1 |
| 1 | 1 | 0 | 1 | 0 | 1 | 1 | 1 |



- Grayscale -> binary mask
- Useful if object of interest's intensity distribution is distinct from background

$$F_{T}[i,j] = \begin{cases} 1 & \text{if } F[i,j] \ge T \\ 0 & \text{otherwise.} \end{cases}$$


$$F_{T}[i,j] = \begin{cases} 1 & \text{if } T_{1} \le F[i,j] \le T_{2} \\ 0 & \text{otherwise.} \end{cases}$$

$$F_T[i,j] = \begin{cases} 1 & \text{if } F[i,j] \in \mathbb{Z} \\ 0 & \text{otherwise.} \end{cases}$$

<u>Example</u>
 http://homepages.inf.ed.ac.uk/rbf/CVonline/LOCAL_COPIES/FITZGIBBON/simplebinary.html

 Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: edge detection



Looking for pixels where gradient is strong.


Given a grayscale image or an intermediate matrix

 threshold to create a binary output.

Example: background subtraction




Looking for pixels that differ significantly from the "empty" background.



fg_pix = find(diff > t);

Given a grayscale image or an intermediate matrix >
threshold to create a binary output.

Example: intensity-based detection




 $fg_pix = find(im < 65);$

Looking for dark pixels

Given a grayscale image or an intermediate matrix


 threshold to create a binary output.

Example: color-based detection




Looking for pixels within a certain hue range.


A nice case: bimodal intensity histograms




Not so nice cases



Two distinct modes




Overlapped modes




Issues

- What to do with "noisy" binary outputs?
 - Holes
 - Extra small fragments

- How to demarcate multiple regions of interest?
 - Count objects
 - Compute further features per object






Morphological operators


- Change the shape of the foreground regions via intersection/union operations between a scanning structuring element and binary image.
- Useful to clean up result from thresholding
- Basic operators are:
 - Dilation
 - Erosion

Dilation

- Expands connected components
- Grow features
- Fill holes



Before dilation




After dilation

Erosion

- Erode connected components
- Shrink features
- Remove bridges, branches, noise




Before erosion



After erosion

Structuring elements

 Masks of varying shapes and sizes used to perform morphology, for example:



 Scan mask across foreground pixels to transform the binary image


>> help strel


Dilation vs. Erosion


At each position:


• **Dilation**: if current pixel is foreground, OR the structuring element with the input image.

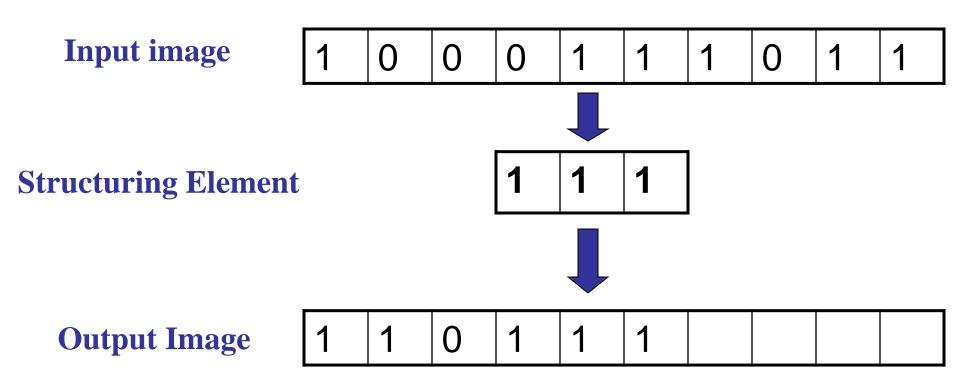
Example for Dilation (1D)




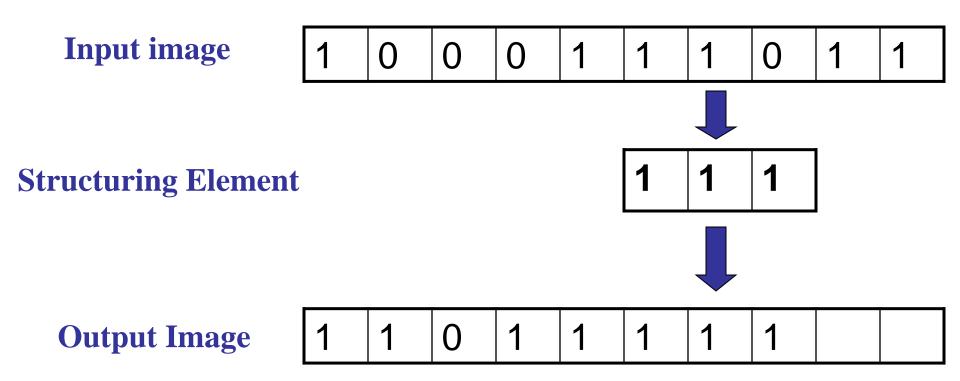



















Note that the object gets bigger and holes are filled.

>> help imdilate

2D example for dilation



| 1_ | 11 | 1 |
|----|----|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |

(a) Binary image B

(b) Structuring element S

| | 1 | 1 | 1 | 1 | 1 | 1 | '1 | 1 |
|---|---|---|---|------|---|----|----|---|
| | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| | | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| L | | 1 | 1 | . 1_ | 1 | 1. | 1 | 1 |
| | | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| | | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| L | | 1 | 1 | 1 | 1 | | | |

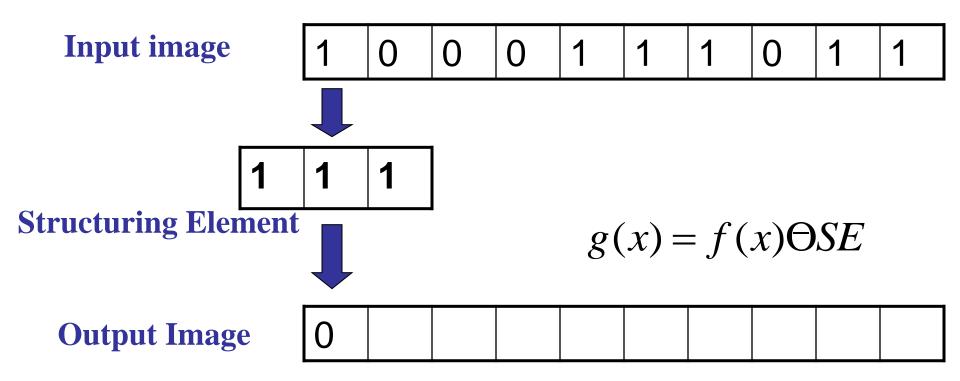
(c) Dilation B ⊕ S

Dilation vs. Erosion

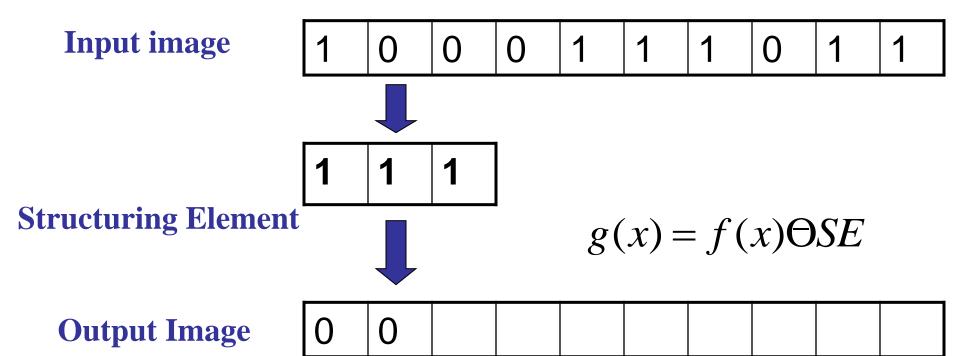
At each position:

- Dilation: if current pixel is foreground, OR the structuring element with the input image.
- **Erosion**: if **every pixel** under the structuring element's nonzero entries is foreground, OR the current pixel with S.


Example for Erosion (1D)




Example for Erosion (1D)





Example for Erosion





Example for Erosion


















Note that the object gets smaller


>> help imerode

2D example for erosion



(a) Binary image B

(b) Structuring element S




(d) Erosion $\mathbf{B} \ominus \mathbf{S}$

Opening

- Erode, then dilate
- Remove small objects, keep original shape




Before opening




After opening

Closing

- Dilate, then erode
- Fill holes, but keep original shape



Before closing



After closing

Applet: http://bigwww.epfl.ch/demo/jmorpho/start.php

Morphology operators on grayscale images


- Dilation and erosion typically performed on binary images.
- If image is grayscale: for dilation take the neighborhood max, for erosion take the min.







dilated





eroded

Issues

- What to do with "noisy" binary outputs?
 - Holes
 - Extra small fragments

- How to demarcate multiple regions of interest?
 - Count objects
 - Compute further features per object






Connected components


Identify distinct regions of "connected pixels"


| 1 | 1 | 0 | 1 | 1 | 1 | 0 | 1 |
|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 1 | 1 | 0 | 0 | 0 | 1 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| 1 | 1 | 1 | 1 | 0 | 1 | 0 | 1 |
| 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 0 | 1 | 0 | 0 | 0 | 1 |
| 1 | 1 | 0 | 1 | 0 | 1 | 1 | 1 |

a) binary image



b) connected components labeling







c) binary image and labeling, expanded for viewing

Connectedness

Defining which pixels are considered neighbors



4-connected



8-connected


Source: Chaitanya Chandra

Connected components

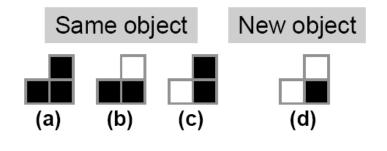
 We'll consider a sequential algorithm that requires only 2 passes over the image.

- Input: binary image
- Output: "label" image, where pixels are numbered per their component

 Note: foreground here is denoted with black pixels.



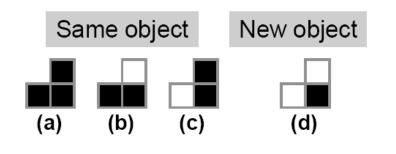
- Labeling a pixel only requires to consider its prior and superior neighbors.
- It depends on the type of connectivity used for foreground (4-connectivity here).




What happens in these cases?




- Labeling a pixel only requires to consider its prior and superior neighbors.
- It depends on the type of connectivity used for foreground (4-connectivity here).

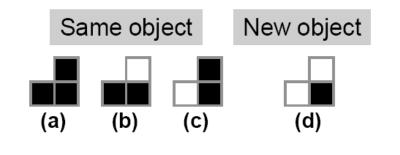


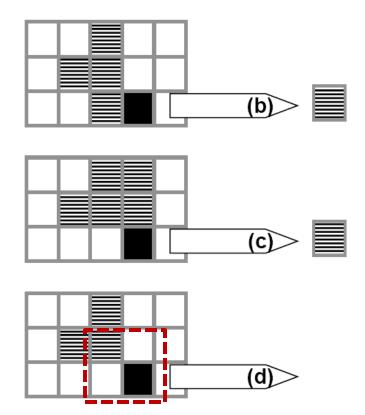


What happens in these cases?

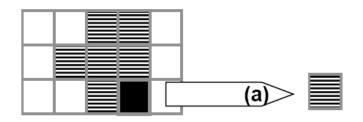


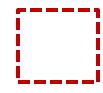
- Labeling a pixel only requires to consider its prior and superior neighbors.
- It depends on the type of connectivity used for foreground (4-connectivity here).





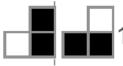
What happens in these cases?





- Process the image from left to right, top to bottom.
 - 1. If the next pixel to process is 1-pixel:


 Already processed



- If only one of its neighbors (<u>superior</u> or <u>left</u>) is 1-pixel, copy its label.
- 2. If both are, and have the same label, copy it.
- 3. If they have different labels:

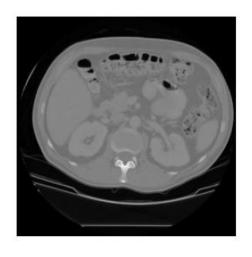
superior? smallest?

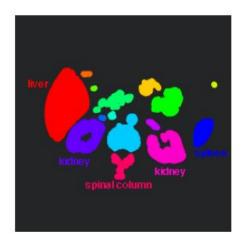
- 1. Copy the label from the prior.
- 2. Reflect the change in the table of equivalences.
- [■]4. Otw, assign a new label.
- 2. More pixels? Go to step 1.



- Re-label with the smallest of equivalent labels.
- Pixels of the same segment always have the same label.

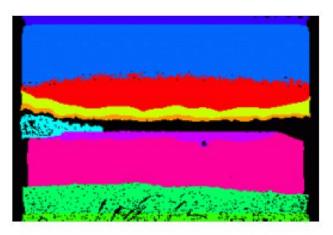
Connected components





connected components of 1's from thresholded image

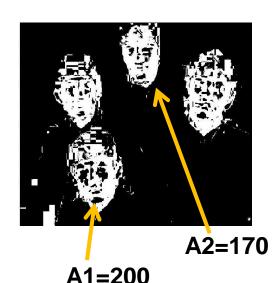


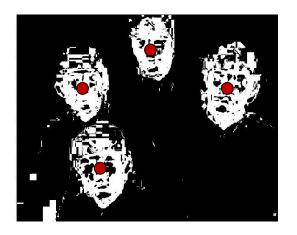


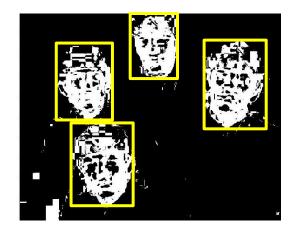
connected components of cluster labels

Region properties

- Given connected components, can compute simple features per blob, such as:
 - Area (num pixels in the region)
 - Centroid (average x and y position of pixels in the region)
 - Bounding box (min and max coordinates)
 - Circularity (ratio of mean dist. to centroid over std)







Circularity

a second measure uses variation off of a circle circularity(2):

$$C_2 = \frac{\mu_R}{\sigma_R}$$

where μ_R and σ_R^2 are the mean and variance of the distance from the centroid of the shape to the boundary pixels (r_k, c_k) .

mean radial distance:

$$\mu_R = \frac{1}{K} \sum_{k=0}^{K-1} \|(r_k, c_k) - (\bar{r}, \bar{c})\|$$

variance of radial distance:

$$\sigma_R^2 = \frac{1}{K} \sum_{k=0}^{K-1} \left[\| (r_k, c_k) - (\bar{r}, \bar{c}) \| - \mu_R \right]^2$$

(r,c) R μ_R

[Haralick]

Binary image analysis: basic steps (recap)

- Convert the image into binary form
 - Thresholding
- Clean up the thresholded image
 - Morphological operators
- Extract separate blobs
 - Connected components
- Describe the blobs with region properties

Matlab

```
• N = hist(Y,M)
L = bwlabel (BW,N);

    STATS = regionprops(L, PROPERTIES) ;

         'Area'
         'Centroid'
         'BoundingBox'
         'Orientation', ...
IM2 = imerode(IM,SE);
IM2 = imdilate(IM,SE);
IM2 = imclose(IM, SE);
IM2 = imopen(IM, SE);
```

Example using binary image analysis: OCR



Digitizing Books One Word at a Time

- → HOME
- → WHAT IS reCAPTCHA

 DIGITIZATION ACCURACY

 WHAT IS A CAPTCHA

 SECURITY
- → GET reCAPTCHA
- → MY ACCOUNT
- → EMAIL PROTECTION
- → RESOURCES




Submit The words above come from scanned books. By typing them, you help to digitize old texts.

reCAPTCHA is a free CAPTCHA service that helps to digitize books, newspapers and old time radio shows. Check out <u>our paper</u> in Science about it (or read more below).

A <u>CAPTCHA</u> is a program that can tell whether its user is a human or a computer. You've probably seen them — colorful images with distorted text at the bottom of Web registration forms. CAPTCHAs are used by many websites to prevent abuse from "bots," or automated programs usually written to generate spam. No computer program can read distorted text as well as humans can, so bots cannot navigate sites protected by CAPTCHAs.

[Luis von Ahn et al. http://recaptcha.net/learnmore.html]

Example using binary image analysis: segmentation of a liver



Example using binary image analysis: Bg subtraction + blob detection



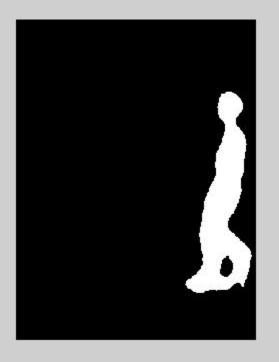








. . .



Example using binary image analysis: Bg subtraction + blob detection



University of Southern California http://iris.usc.edu/~icohen/projects/vace/detection.htm

Binary images

Pros

- Can be fast to compute, easy to store
- Simple processing techniques available
- Lead to some useful compact shape descriptors

Cons

- Hard to get "clean" silhouettes
- Noise common in realistic scenarios
- Can be too coarse of a representation
- Not 3d

Summary

Operations, tools

Derivative filters

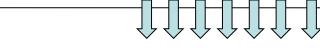
Smoothing, morphology

Thresholding

Connected components

Matched filters

Histograms



 Features, representations

Edges, gradients

Blobs/regions

Local patterns

Textures (next)

Color distributions

Next Class

Features – Corner Detection and SIFT